

LECTURE6 – PARALLEL COMPUTATION PATTERNS (SCAN)



Prefix Sum

A Work-inefficient Scan Kernel

A Work-Efficient Parallel Scan Kernel

More on Parallel Scan



Objective

- To master parallel scan (prefix sum) algorithms
 - Frequently used for parallel work assignment and resource allocation
 - **A key primitive in many parallel algorithms to convert serial computation into parallel computation**
 - A foundational parallel computation pattern
 - Work efficiency in parallel code/algorithms
- Reading –Mark Harris, Parallel Prefix Sum with CUDA
 - http://http.developer.nvidia.com/GPUGems3/gpugems3_ch39.html



Inclusive Scan (Prefix-Sum) Definition

Definition: The **scan** operation takes a binary associative operator \oplus (pronounced as circle plus), and an array of n elements

$$[x_0, x_1, \dots, x_{n-1}],$$

and returns the array

$$[x_0, (x_0 \oplus x_1), \dots, (x_0 \oplus x_1 \oplus \dots \oplus x_{n-1})].$$

Example: If \oplus is addition, then scan operation on the array would return

$$\begin{aligned} &[3 \ 1 \ 7 \ 0 \ 4 \ 1 \ 6 \ 3], \\ &[3 \ 4 \ 11 \ 11 \ 15 \ 16 \ 22 \ 25]. \end{aligned}$$



An Inclusive Scan Application Example

- Assume that we have a 100-inch sandwich to feed 10 people
- We know how much each person wants in inches
 - [3 5 2 7 28 4 3 0 8 1]
- How do we cut the sandwich quickly?
- How much will be left?
- Method 1: cut the sections sequentially: 3 inches first, 5 inches second, 2 inches third, etc.
- Method 2: calculate prefix sum:
 - [3, 8, 10, 17, 45, 49, 52, 52, 60, 61] (39 inches left)



Typical Applications of Scan

- Scan is a simple and useful parallel building block

- Convert recurrences from sequential:

```
for (j=1; j<n; j++)  
    out[j] = out[j-1] + f(j);
```

- Into parallel:

```
forall(j) { temp[j] = f(j) };  
scan(out, temp);
```



Other Applications

- Assigning camping spots
- Assigning Farmer's Market spaces
- Allocating memory to parallel threads
- Allocating memory buffer space for communication channels
- ...



An Inclusive Sequential Addition Scan

Given a sequence $[x_0, x_1, x_2, \dots]$

Calculate output $[y_0, y_1, y_2, \dots]$

Such that $y_0 = x_0$
 $y_1 = x_0 + x_1$
 $y_2 = x_0 + x_1 + x_2$

...

Using a recursive definition

$$y_i = y_{i-1} + x_i$$



A Work Efficient C Implementation

```
y[0] = x[0];  
for (i = 1; i < Max_i; i++) y[i] = y [i-1] + x[i];
```

Computationally efficient:

N additions needed for N elements - $O(N)$!

Only slightly more expensive than sequential reduction.



A Naïve Inclusive Parallel Scan

- Assign one thread to calculate each y element
- Have every thread to add up all x elements needed for the y element

$$y_0 = x_0$$

$$y_1 = x_0 + x_1$$

$$y_2 = x_0 + x_1 + x_2$$

“Parallel programming is easy as long as you do not care about performance.”



Prefix Sum

A Work-inefficient Scan Kernel

A Work-Efficient Parallel Scan Kernel

More on Parallel Scan




Objective

- To learn to write and analyze a high-performance scan kernel
 - Interleaved reduction trees
 - Thread index to data mapping
 - Barrier Synchronization
 - Work efficiency analysis



A Better Parallel Scan Algorithm

1. Read input from device global memory to **shared memory**
2. Iterate **$\log(n)$** times; stride from 1 to $n-1$: **double stride** each iteration




- Active threads *stride* to $n-1$ (n -stride threads)
- Thread j adds elements j and j -stride from shared memory and writes result into element j in shared memory
- Requires barrier synchronization, once before read and once before write



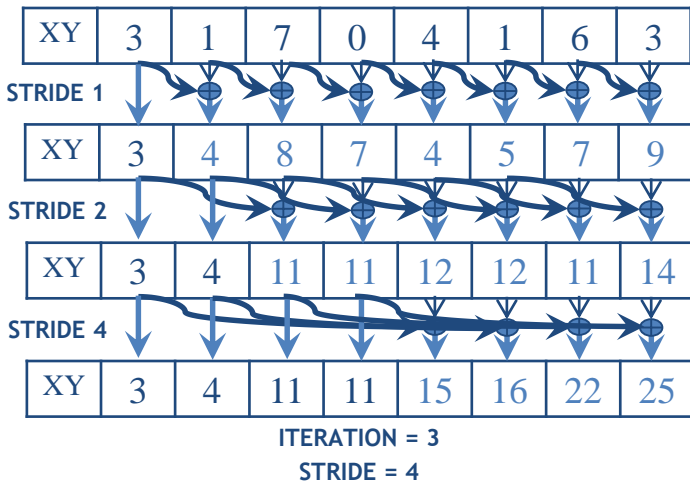
A Better Parallel Scan Algorithm

1. Read input from device to shared memory
2. Iterate $\log(n)$ times; stride from 1 to $n-1$: double stride each iteration.




A Better Parallel Scan Algorithm

1. Read input from device to shared memory
2. Iterate $\log(n)$ times; stride from 1 to $n-1$: double stride each iteration
3. Write output from shared memory to device memory



Handling Dependencies

- During every iteration, each thread can overwrite the input of another thread
 - Barrier synchronization to ensure all inputs have been properly generated
 - All threads secure input operand that can be overwritten by another thread
 - Barrier synchronization is required to ensure that all threads have secured their inputs
 - All threads perform addition and write output



A Work-Inefficient Scan Kernel

```
__global__ void work_inefficient_scan_kernel(float *X, float *Y, int InputSize)
{
    __shared__ float XY[SECTION_SIZE];
    int i = blockIdx.x * blockDim.x + threadIdx.x;
    if (i < InputSize) {XY[threadIdx.x] = X[i];}
    // the code below performs iterative scan on XY
    for (unsigned int stride = 1; stride <= threadIdx.x; stride *= 2) {
        __syncthreads();
        float in1 = XY[threadIdx.x - stride];
        __syncthreads();
        XY[threadIdx.x] += in1;
    }
    __syncthreads();
    If (i < InputSize) {Y[i] = XY[threadIdx.x];}
}
```



Work Efficiency Considerations

- This Scan executes $\log(n)$ parallel iterations
 - The iterations do $(n-1), (n-2), (n-4), \dots (n - n/2)$ adds each
 - Total adds: $n * \log(n) - (n-1) \rightarrow O(n * \log(n))$ work
- This scan algorithm is not work efficient
 - Sequential scan algorithm does n adds
 - A factor of $\log(n)$ can hurt: 10x for 1024 elements!
- A parallel algorithm can be slower than a sequential one when execution resources are saturated from low work efficiency



Prefix Sum

A Work-inefficient Scan Kernel

A Work-Efficient Parallel Scan
Kernel

More on Parallel Scan



Objective

- To learn to write a work-efficient scan kernel
 - Two-phased balanced tree traversal
 - Aggressive re-use of intermediate results
 - Reducing control divergence with more complex thread index to data index mapping



Improving Efficiency

- *Balanced Trees*

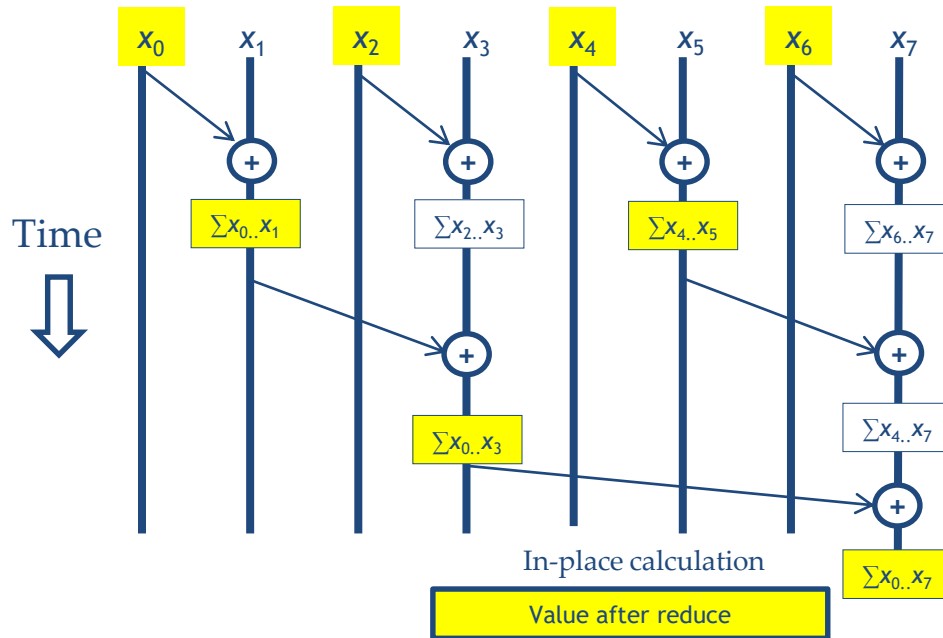
- Form a balanced binary tree on the input data and sweep it to and from the root
- Tree is not an actual data structure, but a concept to determine what each thread does at each step

- For scan:

- Traverse down from leaves to the root building partial sums at internal nodes in the tree
 - The root holds the sum of all leaves
- Traverse back up the tree building the output from the partial sums



Parallel Scan - Reduction Phase



Reduction Phase Kernel Code

```
// XY[2*BLOCK_SIZE] is in shared memory

for (unsigned int stride = 1; stride <= BLOCK_SIZE;
     stride *= 2) {
    int index = (threadIdx.x+1)*stride*2 - 1;
    if(index < 2*BLOCK_SIZE)
        XY[index] += XY[index-stride];
    __syncthreads();
}
```

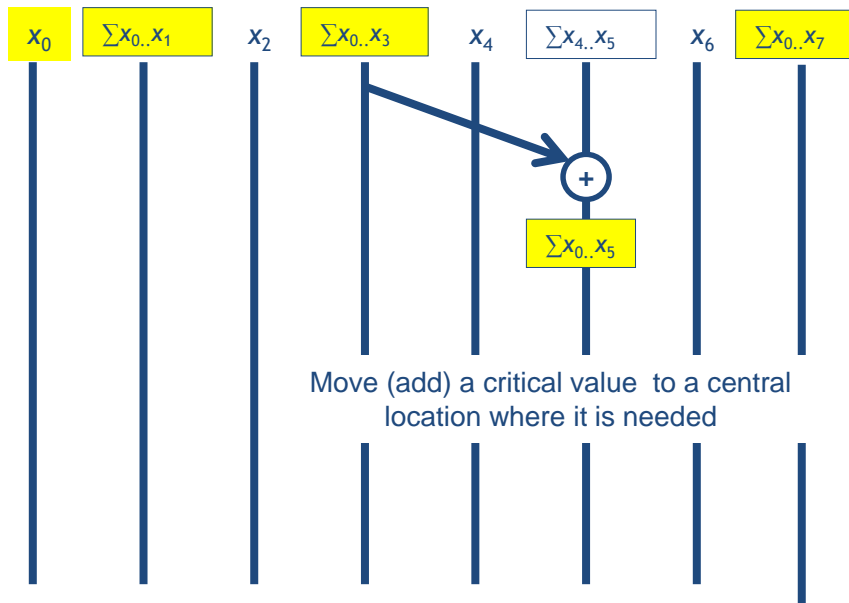
threadIdx.x+1 = 1, 2, 3, 4....

stride = 1,


index = 1, 3, 5, 7, ...




Parallel Scan - Post Reduction Reverse Phase



Parallel Scan - Post Reduction Reverse Phase



Putting it Together



Post Reduction Reverse Phase Kernel Code

```
for (unsigned int stride = BLOCK_SIZE/2; stride > 0; stride /= 2) {  
    __syncthreads();  
    int index = (threadIdx.x+1)*stride*2 - 1;  
    if(index+stride < 2*BLOCK_SIZE) {  
        XY[index + stride] += XY[index];  
    }  
}  
__syncthreads();  
if (i < InputSize) Y[i] = XY[threadIdx.x];
```

First iteration for 16-element section

$\text{threadIdx.x} = 0$

$\text{stride} = \text{BLOCK_SIZE}/2 = 8/2 = 4$

$\text{index} = 8-1 = 7$



Prefix Sum

A Work-inefficient Scan Kernel

A Work-Efficient Parallel Scan Kernel

More on Parallel Scan



Objective

- To learn more about parallel scan
 - Analysis of the work efficient kernel
 - Exclusive scan
 - Handling very large input vectors



Work Analysis of the Work Efficient Kernel

- The work efficient kernel executes $\log(n)$ parallel iterations in the reduction step
 - The iterations do $n/2, n/4, \dots, 1$ adds
 - Total adds: $(n-1) \rightarrow O(n)$ work
- It executes $\log(n)-1$ parallel iterations in the post-reduction reverse step
 - The iterations do $2-1, 4-1, \dots, n/2-1$ adds
 - Total adds: $(n-2) - (\log(n)-1) \rightarrow O(n)$ work
- Both phases perform up to no more than $2 \times (n-1)$ adds
- The total number of adds is no more than twice of that done in the efficient sequential algorithm
 - The benefit of parallelism can easily overcome the $2X$ work when there is sufficient hardware



Some Tradeoffs

- The work efficient scan kernel is normally more desirable
 - Better Energy efficiency
 - Less execution resource requirement
- However, the work inefficient kernel could be better for absolute performance due to its single-phase nature (forward phase only)
 - There is sufficient execution resource

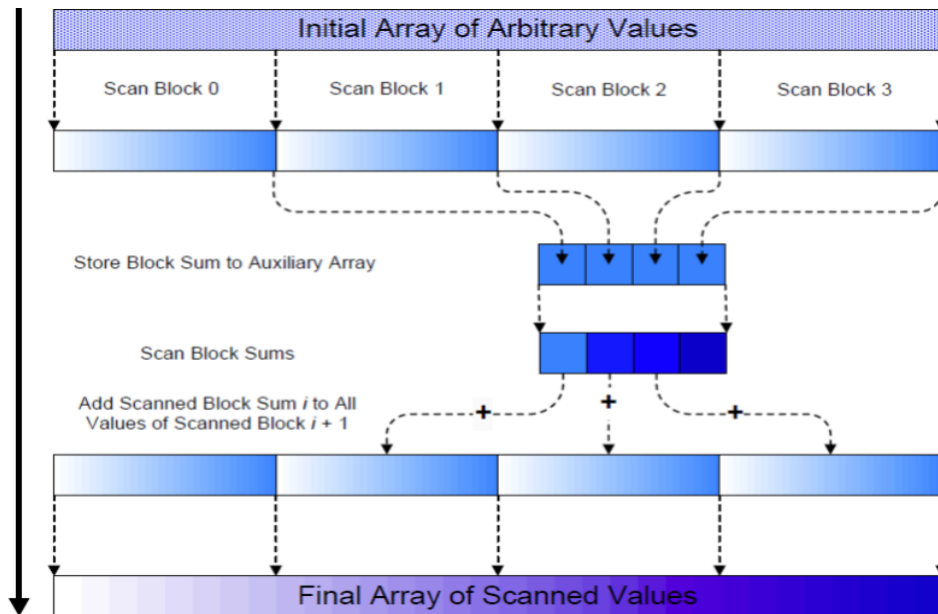


Handling Large Input Vectors

- Build on the work efficient scan kernel
- Have each section of $2 \times \text{blockDim.x}$ elements assigned to a block
 - Perform parallel scan on each section
- Have each block write the sum of its section into a `Sum[]` array indexed by `blockIdx.x`
- Run the scan kernel on the `Sum[]` array
- Add the scanned `Sum[]` array values to all the elements of corresponding sections
- Adaptation of work inefficient kernel is similar.



Overall Flow of Complete Scan



Exclusive Scan Definition

Definition: *The exclusive scan operation takes a binary associative operator \oplus , and an array of n elements*

$$[x_0, x_1, \dots, x_{n-1}]$$

and returns the array

$$[0, x_0, (x_0 \oplus x_1), \dots, (x_0 \oplus x_1 \oplus \dots \oplus x_{n-2})].$$

Example: If \oplus is addition, then the exclusive scan operation on the array $[3 \ 1 \ 7 \ 0 \ 4 \ 1 \ 6 \ 3]$, would return $[0 \ 3 \ 4 \ 11 \ 11 \ 15 \ 16 \ 22]$.



Why Use Exclusive Scan?

- To find the beginning address of allocated buffers
- Inclusive and exclusive scans can be easily derived from each other; it is a matter of convenience

[3 1 7 0 4 1 6 3]

Exclusive [0 3 4 11 11 15 16 22]

Inclusive [3 4 11 11 15 16 22 25]



A Simple Exclusive Scan Kernel

- Adapt an inclusive, work inefficient scan kernel
- Block 0:
 - Thread 0 loads 0 into $XY[0]$
 - Other threads load $X[\text{threadIdx.x}-1]$ into $XY[\text{threadIdx.x}]$
- All other blocks:
 - All thread load $X[\text{blockIdx.x}*\text{blockDim.x}+\text{threadIdx.x}-1]$ into $XY[\text{threadIdx.x}]$
- Similar adaption for work efficient scan kernel but ensure that each thread loads two elements
 - Only one zero should be loaded
 - All elements should be shifted to the right by only one position

Read the Harris article (Parallel Prefix Sum with CUDA) for a more intellectually interesting approach to exclusive scan kernel implementation.

